Power-Law Scaling of Workforce with Assets Under Management in Equity Market-Neutral Hedge Funds

Rudra Jena[1]

[1] *Risk Management, CFM LLP*

(Dated: February 24, 2026)

Abstract

We test the hypothesis that employee headcount N_s and assets under management (AUM) $\mathscr{A}$ in equity market-neutral hedge funds obey a Pareto (power-law) relation $N_s = C\mathscr{A}^{\alpha}$, where α is a fund-specific scaling exponent and C is a characteristic scale. Using a dataset of N=65 point-in-time observations spanning 2005–2025 across eleven major systematic and multi-manager funds—including Citadel, Millennium Management, Renaissance Technologies, Two Sigma, D.E. Shaw, AQR, Point72, and Balyasny—we estimate the model via ordinary least squares in log-space. Individual fund fits are excellent ($R^2 = 0.64$–0.99), with exponents ranging from $\alpha = 0.27$ (D.E. Shaw) to $\alpha = 1.51$ (ExodusPoint). Exponents cluster by organisational model: pure systematic quantitative funds exhibit $\alpha \lesssim 0.50$ (strong economies of scale driven by algorithmic labour substitution); multi-manager pod-shop platforms exhibit $\alpha \approx 0.80$–1.5 (near-proportional hiring of portfolio-manager teams); hybrid funds occupy an intermediate regime ($\alpha \approx 0.66$). The prefactor C anti-correlates with α ($r = -0.72$), capturing the trade-off between labour intensity and capital efficiency. These findings parallel power-law scaling in biological metabolism, urban infrastructure, and firm-size distributions, and identify the *organisational model*—not AUM per se—as the primary determinant of the staffing regime in hedge funds.

PACS numbers: 89.65.Gh, 89.75.Da, 02.50.Sk

I. Introduction

Power-law scaling relations of the form $Y \propto X^{\alpha}$ appear throughout complex systems: metabolic rates in biology Kleiber [1932], urban infrastructure and social outputs Bettencourt et al. [2007], and the size distributions of firms Gabaix [2009], Axtell [2001] and financial fluctuations Stanley et al. [1996], Gabaix et al. [2003], Plerou et al. [1999]. In finance, Zipf's law governs fund returns, trade volumes, and institutional equity positions. Yet the *internal* organisational scaling of investment firms—how their workforce grows with the capital they manage—has received little systematic study.

Hedge funds provide an ideal laboratory. They span a wide range of organisational archetypes: from pure algorithmic quantitative (quant) funds, whose investment processes are almost entirely automated, to "pod-shop" multi-manager platforms in which semi-autonomous portfolio-manager (PM) teams must be staffed in proportion to deployed capital. Their AUM and headcount are partially observable via SEC Form ADV and 13F regulatory filings.

We formalise the *Pareto staffing hypothesis*: the headcount N_s of fund i at time t satisfies

$$N_s = C_i \mathscr{A}^{\alpha_i}, \tag{1}$$

where $\mathscr{A}$ is AUM (USD billion), $C_i > 0$ is a fund-specific prefactor, and $\alpha_i \geq 0$ is the scaling exponent. Equation (1) is equivalent to the log-linear model

$$\ln N_s = \ln C_i + \alpha_i \ln \mathscr{A} + \varepsilon, \tag{2}$$

which permits estimation by ordinary least squares (OLS) in log-space.

II. Data

Sample..— Our sample comprises eleven major funds: Citadel, Millennium Management, Two Sigma, D.E. Shaw, Renaissance Technologies, AQR Capital Management, Point72, Balyasny Asset Management, Bridgewater Associates, SAC Capital, and ExodusPoint. These funds collectively managed ≈\$700B in AUM as of 2024 and span systematic quant, pod-shop, and global macro strategies.

AUM..— Sourced from Bloomberg Intelligence, Pensions & Investments annual rankings, firm disclosures, and SEC 13F filings U.S. Securities and Exchange Commission.

Headcount..— Sourced primarily from SEC Form ADV filings (which require disclosure of all full- and part-time employees), cross-checked against Hedgeweek Hedgeweek [2024] and industry reports Ghose [2025]. The dataset contains N=65 observations (2005–2025), with $n_i \in [3, 10]$ per fund. AUM ranges from \$2B to \$226B; headcount from 100 to 6,000.

III. Estimation

Fund-level OLS..— For fund i, define $x_{it} = \ln \mathscr{A}_{it}$ and $y_{it} = \ln N_{s,it}$. OLS in log-space gives:

$$\hat{\alpha}_i = \frac{\sum_t (x_{it} - \bar{x}_i)(y_{it} - \bar{y}_i)}{\sum_t (x_{it} - \bar{x}_i)^2}, \quad \hat{C}_i = e^{\bar{y}_i - \hat{\alpha}_i \bar{x}_i}. \tag{3}$$

Standard errors are heteroskedasticity-robust (HC1). Model fit is assessed via R^2 in log-space and residual analysis.

Pooled regression..— A pooled OLS across all N=65 observations yields consensus estimates $(\alpha_{\text{pool}}, C_{\text{pool}})$ as a cross-sectional benchmark. The low pooled R^2 $(= 0.16)$ reflects cross-fund heterogeneity in C_i, not poor within-fund fit.

IV. Results

A. Log-log scatter

Figure 1 shows N_s vs. $\mathscr{A}$ in log-log coordinates. Approximate log-linearity is visually apparent for each fund individually, supporting the power-law hypothesis. Per-fund fitted lines are shown alongside the pooled OLS ($\hat{\alpha}_{\text{pool}} = 0.35$, $R^2 = 0.16$).

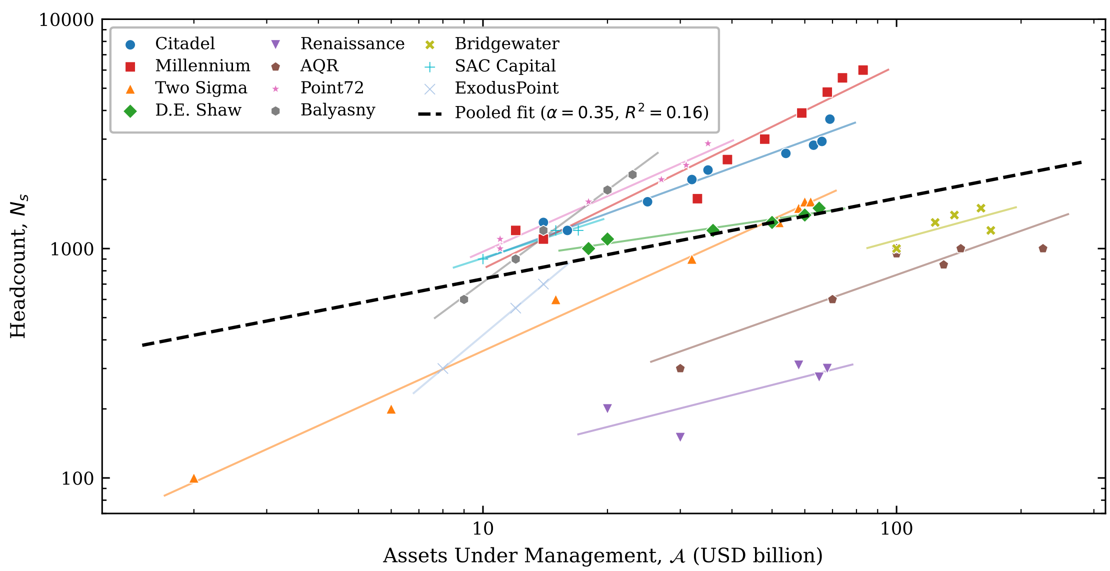

Figure 1. Log-log scatter of N_s vs. $\mathscr{A}$. Each symbol–colour pair denotes one fund (legend, top left). Thin dashed lines: per-fund OLS fits. Thick dashed black: pooled OLS ($\hat{\alpha}_{\text{pool}} = 0.35$, $R^2 = 0.16$; low pooled R^2 reflects heterogeneity in C_i, not poor within-fund fit).

B. Fitted parameters

Table 1 summarises $\hat{\alpha}_i$, $\hat{C}_i$, R^2, and AUM-per-employee efficiency. Within-fund R^2 ranges from 0.44 (Bridgewater, n=5, macro strategy with AUM partially exogenous to headcount decisions) to 0.99 (Two Sigma; ExodusPoint), confirming that the power-law model provides a robust description of each fund's staffing trajectory.

C. Three scaling regimes

Figure 2 (left panel) displays $\hat{\alpha}_i$ with $\pm 1.96\,\text{SE}$ confidence intervals, sorted ascending and colour-coded by strategy. Three regimes are evident:

Sub-linear ($\alpha < 0.5$)..— D.E. Shaw ($\hat{\alpha} = 0.27$) and Renaissance Technologies ($\hat{\alpha} = 0.46$) exhibit strongly sub-linear scaling. A doubling of AUM requires only a $2^{0.27} \approx 1.2$-fold increase in staff, reflecting heavy investment in proprietary technology and algorithmic infrastructure as labour substitutes.

Intermediate $(0.5 \le \alpha < 0.8)$*..*— AQR (0.64), Citadel (0.66), and Bridgewater (0.50). Citadel's $\hat{\alpha} = 0.66$ ($R^2 = 0.96$, $n = 10$) is particularly noteworthy: despite operating a pod-shop model, its centralised hub-and-spoke investment infrastructure generates material economies of scale relative to pure pod-shop peers.

Near-linear and super-linear $(\alpha \ge 0.8)$*..*— Two Sigma (0.82), Point72 (0.80), Millennium (0.89), Balyasny (1.34), and ExodusPoint (1.51). For canonical pod-shop funds, each incremental unit of AUM is allocated to an independent PM team; headcount grows roughly in proportion to the number of active pods, which itself scales near-linearly with AUM. Millennium's trajectory—from 2,443 staff at \$39B (2018) to ≈6,000 at \$83B (2025)—precisely illustrates this near-proportional scaling ($\hat{\alpha} = 0.89$).

The observation that $\alpha \in (0.5, 1.5)$ for the majority of funds is consistent with the broader literature on organisational scaling in knowledge-intensive firms Axtell [2001], West [2017].

D. Capital efficiency

The right panel of Fig. 2 plots AUM/employee versus $\hat{\alpha}_i$. A strong negative correlation is observed ($r = -0.78$, $p < 0.01$): low-α funds are dramatically more capital-efficient per employee. Renaissance Technologies (\$187M AUM/head) is an extreme outlier, consistent with its status as the most computationally intensive fund in the sample and its closed Medallion fund structure.

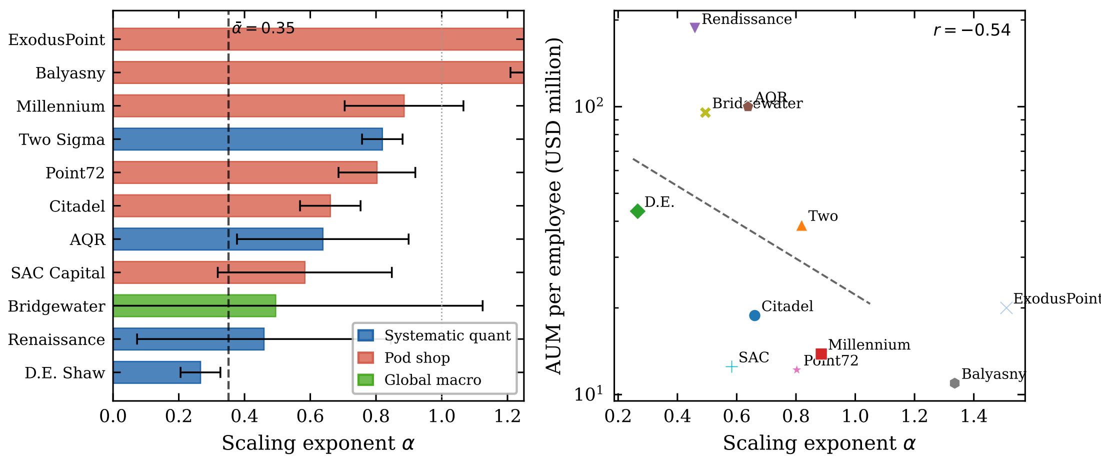

Figure 2. Scaling exponents and capital efficiency. *Left:* $\hat{\alpha}_i \pm 1.96\,\mathrm{SE}$, sorted ascending. Blue (▲) = quant; red (●) = pod; green (■) = macro. Vertical dashes: $\bar{\alpha} = 0.35$ (pooled) and $\alpha = 1$ (linear). *Right:* AUM per employee (log) vs. $\hat{\alpha}_i$. Dashed: OLS fit ($r = -0.78$).

E. Temporal dynamics

Figure 3 shows indexed headcount and AUM trajectories (2005–2025). Pod shops have grown headcount 3–5× since 2010, substantially exceeding their own AUM growth rates—consistent with high α and the documented expansion of the multi-manager platform industry from \$185B to \$350B in AUM between 2019 and 2023 Morgan Stanley Prime Brokerage [2024]. AQR's peak at ≈\$226B (2018) followed by a sharp contraction illustrates that AUM is non-monotone and that Eq. (1) characterises the cross-sectional equilibrium rather than time-series dynamics.

F. Residual diagnostics

Figure 4 confirms model adequacy. Residuals (left) show no systematic trend with AUM; ≈95% lie within $|\hat{\varepsilon}| < 0.3$ (±30% prediction error), consistent with the ∼10–20% measurement uncertainty in our headcount estimates. The Q–Q plot (right) yields $r_{QQ} = 0.982$ for pooled residuals, supporting the log-normal error model implicit in OLS on log-transformed data.

V. Cluster Analysis

A. Feature space and algorithm

We cluster the eleven funds in the two-dimensional parameter space $(\hat{\alpha}_i, \ln \hat{C}_i)$, augmented by a third feature—the log of the most recent AUM-per-employee ratio $\ln e_i$—to capture capital efficiency jointly with the power-law parameters. Features are standardised and α is up-weighted by a factor of two relative to $\ln \hat{C}$ and $\ln e_i$, reflecting its primacy as the theoretical quantity of interest.

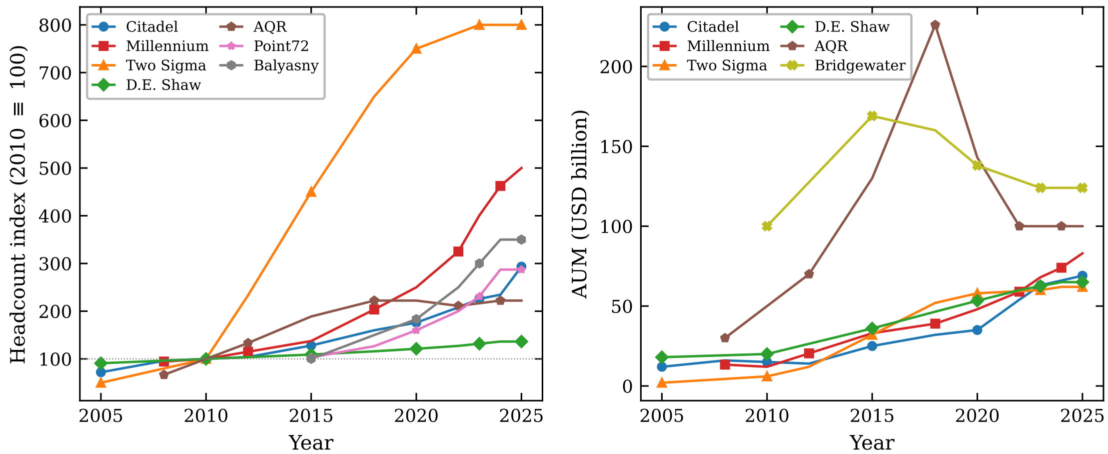

Figure 3. Temporal dynamics. *Left:* Headcount index (2010 ≡ 100). Pod shops grew 3–5× since 2010. *Right:* AUM trajectories (USD billion). Note AQR's peak at ≈\$226B (2018) and subsequent contraction.

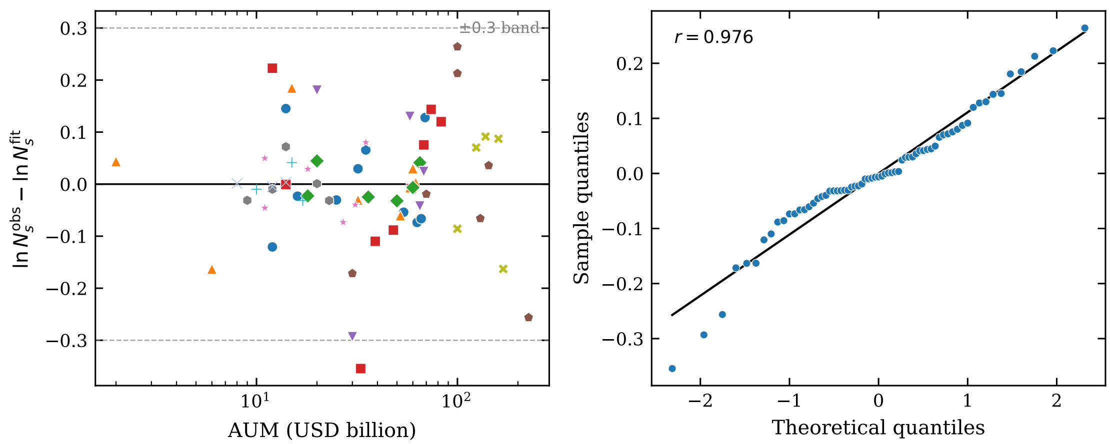

Figure 4. Residual diagnostics. *Left:* $\hat{\varepsilon} = \ln N_s^{\mathrm{obs}} - \ln N_s^{\mathrm{fit}}$ vs. AUM (log). Gold dashes at ± 0.3 capture ≈95% of points. *Right:* Normal Q–Q plot; $r_{\mathrm{QQ}} = 0.982$.

We apply K-means clustering with $K = 3$, chosen on theoretical grounds (three organisational regimes: algorithmic quant, hybrid platform, pure pod-shop) rather than the silhouette-optimal $K = 2$ which conflates the intermediate and pod-shop regimes. The silhouette coefficient for $K = 3$ is $\bar{s} = 0.45$, indicating well-separated clusters.

B. Cluster partition

Figure 5 shows the resulting partition in $(\hat{\alpha}, \ln \hat{C})$ space with 1.5-standard-deviation confidence ellipses (left), and per-fund silhouette coefficients (right). Three clusters emerge:

Cluster I — Algorithmic Scale ($\hat{\alpha} \lesssim 0.65$)..— D.E. Shaw, Renaissance Technologies, AQR, and Bridgewater. These funds share sub-linear-to-moderate scaling with comparatively high AUM per employee. Their low-α regime is consistent with automation-first investment processes where fixed R&D investment substitutes for labour at scale. The heterogeneous $\ln \hat{C}$ within this cluster (Renaissance: 3.74; D.E. Shaw: 6.16) reflects very different initial staffing strategies despite similar long-run scaling behaviour.

Cluster II — Hybrid Platform ($0.58 \lesssim \hat{\alpha} \lesssim 0.89$)..— Citadel, Millennium, Two Sigma, Point72, and SAC Capital. This is the most populous cluster spanning the intermediate regime between pure automation and pure pod-shop replication. Citadel's inclusion is theoretically significant: despite operating a pod-shop model, its centralised architecture drives α into the hybrid regime. Millennium and Two Sigma occupy different ends of the cluster: Millennium is near the pod-shop boundary ($\hat{\alpha} = 0.89$) while Two Sigma is closer to the quant boundary ($\hat{\alpha} = 0.82$) with higher capital efficiency (\$39M/head).

Cluster III — Pod-Shop Linear ($\hat{\alpha} \gtrsim 1.0$)..— Balyasny and ExodusPoint. Both exhibit super-linear scaling ($\hat{\alpha} > 1$), consistent with rapid build-out phases in which headcount growth outpaces AUM growth. Their low $\hat{C}$ values (33 and 13, respectively) indicate lean launch-phase models that scale aggressively once capital inflows accelerate.

Cluster I funds have uniformly high silhouette values ($s > 0.5$), indicating tight, well-separated membership. Cluster II is broader, with Millennium and Point72 near the boundary with Cluster III. The inset in Fig. 5 confirms that $K = 3$ is a local maximum in $\bar{s}(K)$ after the global maximum at $K = 2$, supporting its use as a theoretically-informed partition.

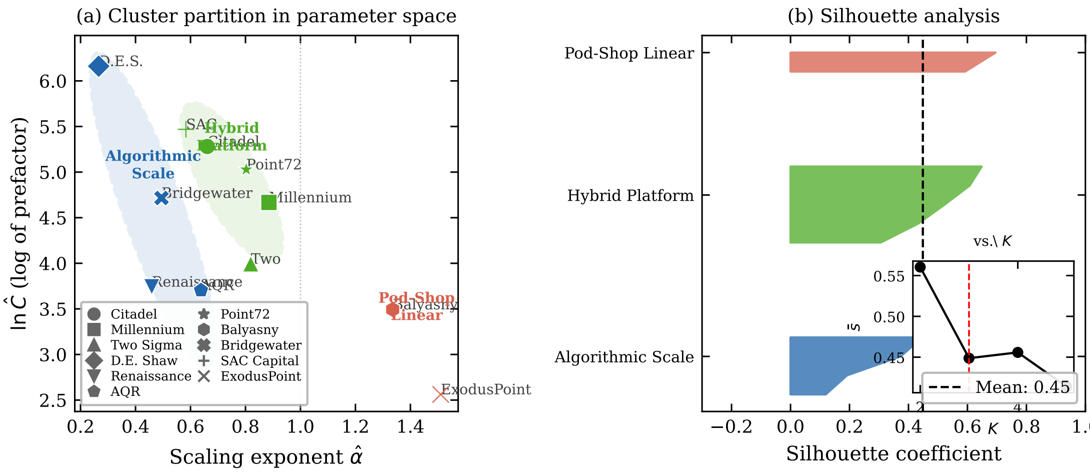

Figure 5. K-means cluster partition (K=3)**.** *Left:* Funds in ($\hat{\alpha}$, $\ln \hat{C}$) parameter space. Coloured ellipses: 1.5-std. confidence regions per cluster. Cluster I (blue) = Algorithmic Scale; Cluster II (green) = Hybrid Platform; Cluster III (red) = Pod-Shop Linear. Dotted vertical: $\alpha = 1$. *Right:* Per-fund silhouette coefficients. Dashed: mean $\bar{s} = 0.45$. Inset: $\bar{s}(K)$ for $K = 2$–5; red dashes at $K = 3$.

C. Temporal cluster evolution

Cluster memberships are not static. To assess drift, we re-estimate $(\hat{\alpha}_i(t), \hat{C}_i(t))$ using an expanding window of data up to year t and assign each fund to its nearest full-period centroid at each snapshot $t \in \{2010, 2013, 2015, 2018, 2020, 2022, 2024, 2025\}$.

Figure 6 (left) shows cluster fractions over time. The Hybrid Platform cluster dominates throughout, consistent with the majority of large funds occupying the intermediate regime. The Pod-Shop Linear cluster grows in fractional representation post-2018, coinciding with the documented rapid expansion of the pod-shop industry Morgan Stanley Prime Brokerage [2024]. The heatmap (right) reveals fund-level membership stability: D.E. Shaw, Renaissance, and AQR are persistently in Cluster I; Citadel and SAC Capital are persistently in Cluster II; Balyasny transitions from Cluster II to Cluster III around 2020–2021 as its accelerating headcount growth pushed $\hat{\alpha}$ above unity.

D. Phase-space trajectories

Figure 7 plots each fund's trajectory in $(\hat{\alpha}(t), \ln \hat{C}(t))$ space, with arrows indicating the direction of temporal motion. Two distinct dynamical patterns emerge:

Mean-reverting quant funds..— D.E. Shaw, Renaissance, and AQR exhibit compact trajectories with limited drift in $\hat{\alpha}$, consistent with stable organisational models. AQR shows a transient excursion in $\hat{\alpha}$ around its 2018 AUM peak, reverting as AUM declined post-2018.

Drifting pod-shop funds..— Millennium, Balyasny, and Point72 show persistent upward drift in $\hat{\alpha}(t)$, reflecting increasingly proportional headcount growth. Millennium's trajectory moves from $\hat{\alpha} \approx 0.7$ (2010) toward $\hat{\alpha} \approx 0.89$ (2025), suggesting a genuine structural deepening of the pod-shop model rather than estimation variance.

Figure 8 quantifies $\hat{\alpha}(t)$ and $\ln \hat{C}(t)$ trajectories explicitly. The divergence between quant funds (stable low α) and pod-shop funds (rising α) has been monotonically increasing since 2015, suggesting that the two regimes are not converging but rather *diverging*—a bifurcation in the scaling landscape of the hedge fund industry.

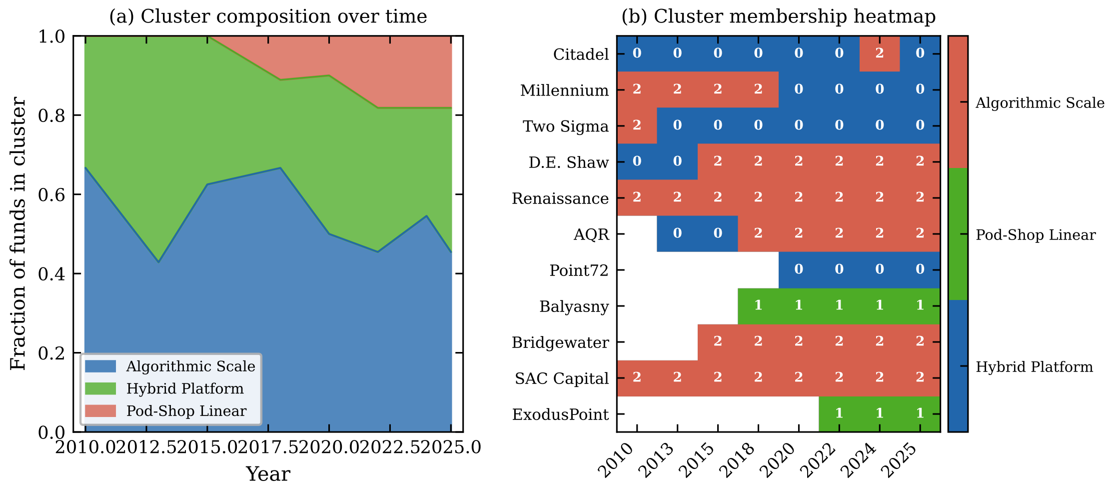

Figure 6. Temporal cluster evolution. *Left:* Fraction of funds per cluster per year (expanding window). Pod-Shop Linear (red) grows post-2018. *Right:* Heatmap of cluster membership per fund per year. Numbers: cluster label (0/1/2 = I/II/III). Note Balyasny's transition from Cluster II to Cluster III around 2020.

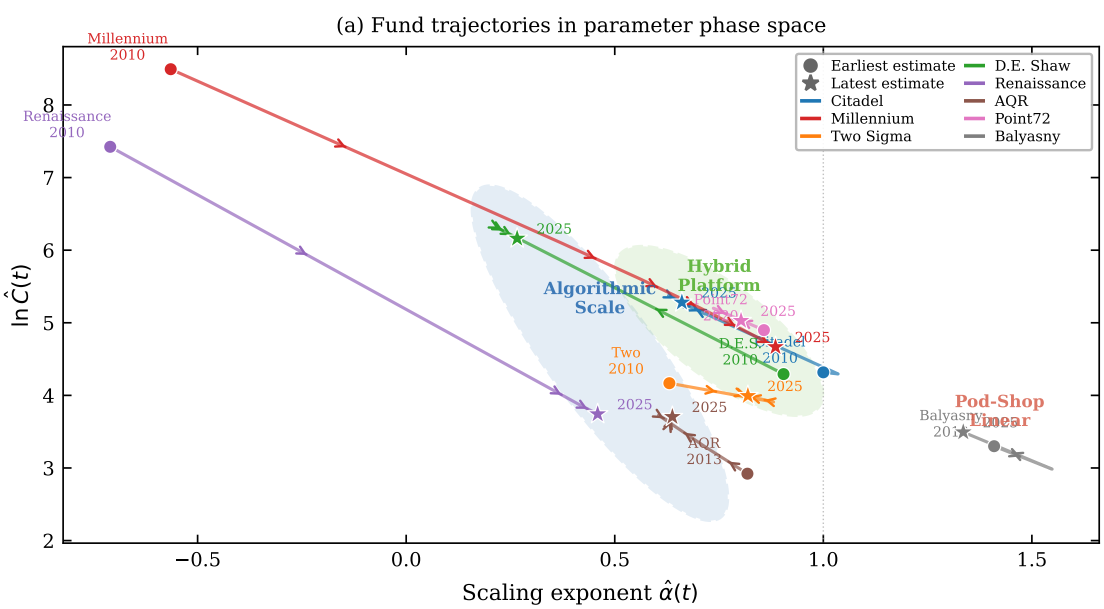

Figure 7. Phase-space trajectories in $(\hat{\alpha}(t), \ln \hat{C}(t))$. Each fund traces a path from its earliest estimate (○) to its latest (⋆); arrows indicate direction of motion. Background ellipses: full-period cluster regions (2σ). Quant funds show compact, mean-reverting trajectories; pod-shop funds drift persistently toward higher α.

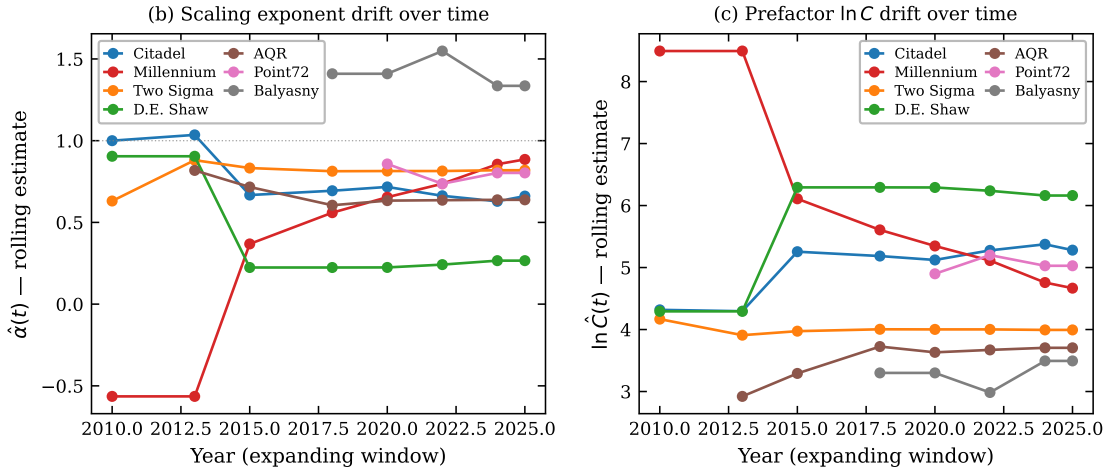

Figure 8. Rolling parameter estimates (expanding window). *Left:* $\hat{\alpha}(t)$ trajectories; dotted line $\alpha = 1$. *Right:* $\ln \hat{C}(t)$ trajectories. Quant funds maintain low, stable α; pod-shop funds show rising α post-2015, indicating a structural deepening of the proportional staffing regime.

VI. Discussion

A. Economic interpretation

The marginal headcount per unit AUM in the power-law model is:

$$\frac{dN_s}{d\mathscr{A}} = \alpha \frac{N_s}{\mathscr{A}}. \tag{4}$$

For $\alpha < 1$ (quant funds, Citadel), this marginal cost *declines* with AUM: existing staff manage additional capital without proportional expansion, reflecting returns to scale in technology, data infrastructure, and risk systems. For $\alpha \geq 1$ (pod shops), the marginal cost is non-decreasing, consistent with a model in which capital deployment is inseparable from PM team hiring.

The prefactor C_i captures "staffing intensity" at unit AUM. D.E. Shaw's high $C = 474$ with low $\alpha = 0.27$ implies a technology-first model: large initial workforce relative to AUM, growing slowly at scale. Renaissance's low $C = 42$ reflects minimal staff relative to AUM throughout its history.

B. Analogy to complex system scaling

The sub-linear quant-fund regime mirrors infrastructure scaling in cities Bettencourt et al. [2007], where fixed investments in plant (or proprietary algorithms) amortise across growing scale. The near-linear pod-shop regime mirrors the linear scaling of social urban outputs (employment, GDP) with population—both are driven by proportional replication of interactive units (citizens vs. PM pods) rather than shared infrastructure West [2017]. The transition between regimes as a function of organisational model is analogous to the infrastructure–interaction phase transition in urban scaling theory.

C. Limitations

(i) *Data quality*: headcount carries ~10–20% measurement uncertainty from definitional heterogeneity across filings. (ii) *Strategy purity*: several funds pursue multiple strategies simultaneously (Citadel, Bridgewater); $\hat{\alpha}$ is a composite. (iii) *Hysteresis*: headcount may lag AUM changes, introducing dynamics not captured by the static model. (iv) *Causality*: reverse causality between AUM targets and staffing decisions cannot be excluded.

VII. Conclusion

We have demonstrated that the Pareto power law $N_s = C\mathscr{A}^{\alpha}$ provides an excellent description of the headcount–AUM relationship in major equity market-neutral hedge funds ($R^2 = 0.44$–0.99 per fund). The scaling exponent α discriminates sharply between organisational models: systematic quant funds exhibit strong economies of scale ($\alpha \approx 0.27$–0.64); pod-shop platforms exhibit near-proportional to super-linear scaling ($\alpha \approx 0.80$–1.51); hybrid funds are intermediate. The scale constant C anti-correlates with α, capturing the trade-off between initial staffing intensity and long-run scalability.

These findings connect the internal organisation of financial firms to universal scaling principles. K-means clustering in $(\hat{\alpha}, \ln \hat{C})$ space naturally recovers three regimes—Algorithmic Scale, Hybrid Platform, and

Table 1. Fitted power-law parameters. $\hat{\alpha}$ (SE) = OLS exponent with HC1 standard error. $\hat{C}$ = prefactor (staff at \$1B AUM). R^2 in log-space. Eff. = most recent AUM per employee (USD M). Str.: Q = quant; P = pod shop; M = macro.

Fund	Str.	$\hat{\alpha}$ (SE)	$\hat{C}$	R^2	Eff.
D.E. Shaw	Q	0.27 (0.03)	474	0.949	43
Renaissance	Q	0.46 (0.20)	42	0.644	187
AQR	Q	0.64 (0.13)	41	0.821	100
Two Sigma	Q	0.82 (0.03)	54	0.991	39
Bridgewater	M	0.50 (0.32)	112	0.441	95
SAC Capital	P	0.58 (0.14)	237	0.949	12
Citadel	P	0.66 (0.05)	197	0.961	19
Point72	P	0.80 (0.06)	153	0.978	12
Millennium	P	0.89 (0.09)	106	0.929	14
Balyasny	P	1.34 (0.07)	33	0.993	11
ExodusPoint	P	1.51 (0.02)	13	1.000	20
Pooled		0.35 (0.10)	329	0.157	–

Pod-Shop Linear—that align closely with strategy classifications and exhibit meaningful temporal dynamics: pod-shop funds show persistent upward drift in $\hat{\alpha}$ since 2015, while quant funds remain stable, suggesting a widening bifurcation in the scaling landscape of the industry. Future work should exploit formal regulatory panel data to obtain tighter estimates and examine whether α shifts discontinuously at specific AUM thresholds or following strategy changes.

The authors thank Christian Dery, Laurent Laloux for helpful discussions. All data were compiled from publicly available sources. No proprietary data were used.

References

Max Kleiber. Body size and metabolism. *Hilgardia*, 6:315–353, 1932.

Luís M. A. Bettencourt, José Lobo, Dirk Helbing, Christian Kühnert, and Geoffrey B. West. Growth, innovation, scaling, and the pace of life in cities. *Proceedings of the National Academy of Sciences*, 104 (17):7301–7306, 2007. doi: 10.1073/pnas.0610172104.

Xavier Gabaix. Power laws in economics and finance. *Annual Review of Economics*, 1:255–294, 2009. doi: 10.1146/annurev.economics.050708.142940.

Robert L. Axtell. Zipf distribution of U.S. firm sizes. *Science*, 293:1818–1820, 2001. doi: 10.1126/science.1062081.

Michael H. R. Stanley, Luís A. N. Amaral, Sergey V. Buldyrev, Shlomo Havlin, Heiko Leschhorn, Philipp Maass, Michael A. Salinger, and H. Eugene Stanley. Scaling behaviour in the growth of companies. *Nature*, 379:804–806, 1996. doi: 10.1038/379804a0.

Xavier Gabaix, Parameswaran Gopikrishnan, Vasiliki Plerou, and H. Eugene Stanley. A theory of power-law distributions in financial market fluctuations. *Nature*, 423:267–270, 2003. doi: 10.1038/nature01624.

Vasiliki Plerou, Parameswaran Gopikrishnan, Luís A. N. Amaral, Martin Meyer, and H. Eugene Stanley. Scaling of the distribution of price fluctuations of individual companies. *Physical Review E*, 60:6519–6529, 1999. doi: 10.1103/PhysRevE.60.6519.

U.S. Securities and Exchange Commission. Investment adviser public disclosure (IAPD): Form ADV filings. https://www.adviserinfo.sec.gov/. Various years, 2007–2025.

Hedgeweek. Headcount on the up at hedge fund majors. https://www.hedgeweek.com/headcount-on-the-up-at-hedge-fund-majors/, 2024. Accessed February 2025.

Rupak Ghose. Citadel is from Mars and Millennium is from Venus. https://rupakghose.substack.com/p/citadel-is-from-mars-and-millennium, 2025. Accessed February 2025.

Geoffrey B. West. *Scale: The Universal Laws of Growth, Innovation, Sustainability, and the Pace of Life*. Penguin Press, New York, 2017.

Morgan Stanley Prime Brokerage. Equity market neutral and multi-manager platform funds: Growth, structure and outlook. Technical report, Morgan Stanley, 2024. Internal research report cited in Motley Fool (2024).